

art-attack-analysis-and-prediction

October 27, 2024

I've been working with the Heart Attack Analysis & Prediction dataset I found on Kaggle. Here's the link to the dataset: [Heart Attack Analysis & Prediction Dataset](#).

```
[3]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, Ridge, Lasso, ElasticNet
from sklearn.preprocessing import PolynomialFeatures, StandardScaler, \
    OneHotEncoder
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.pipeline import Pipeline, make_pipeline
from sklearn.utils import shuffle
from sklearn.model_selection import cross_val_score, cross_val_predict, \
    cross_validate
from sklearn.linear_model import SGDRegressor
from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer
```

Displaying the statistical values for each of the attributes, along with visualizations (e.g., histogram) of the distributions for each attribute.

```
[4]: # Load the data
df = pd.read_csv("heart.csv")
```

Statistical descriptions

```
[5]: # Display summary statistics
statistics = df.describe()
print(statistics)
```

	age	sex	cp	trtbps	chol	fbs	\
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	
mean	54.366337	0.683168	0.966997	131.623762	246.264026	0.148515	
std	9.082101	0.466011	1.032052	17.538143	51.830751	0.356198	
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	
25%	47.500000	0.000000	0.000000	120.000000	211.000000	0.000000	
50%	55.000000	1.000000	1.000000	130.000000	240.000000	0.000000	

75%	61.000000	1.000000	2.000000	140.000000	274.500000	0.000000
max	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000

	restecg	thalachh	exng	oldpeak	slp	caa \
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000
mean	0.528053	149.646865	0.326733	1.039604	1.399340	0.729373
std	0.525860	22.905161	0.469794	1.161075	0.616226	1.022606
min	0.000000	71.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	133.500000	0.000000	0.000000	1.000000	0.000000
50%	1.000000	153.000000	0.000000	0.800000	1.000000	0.000000
75%	1.000000	166.000000	1.000000	1.600000	2.000000	1.000000
max	2.000000	202.000000	1.000000	6.200000	2.000000	4.000000

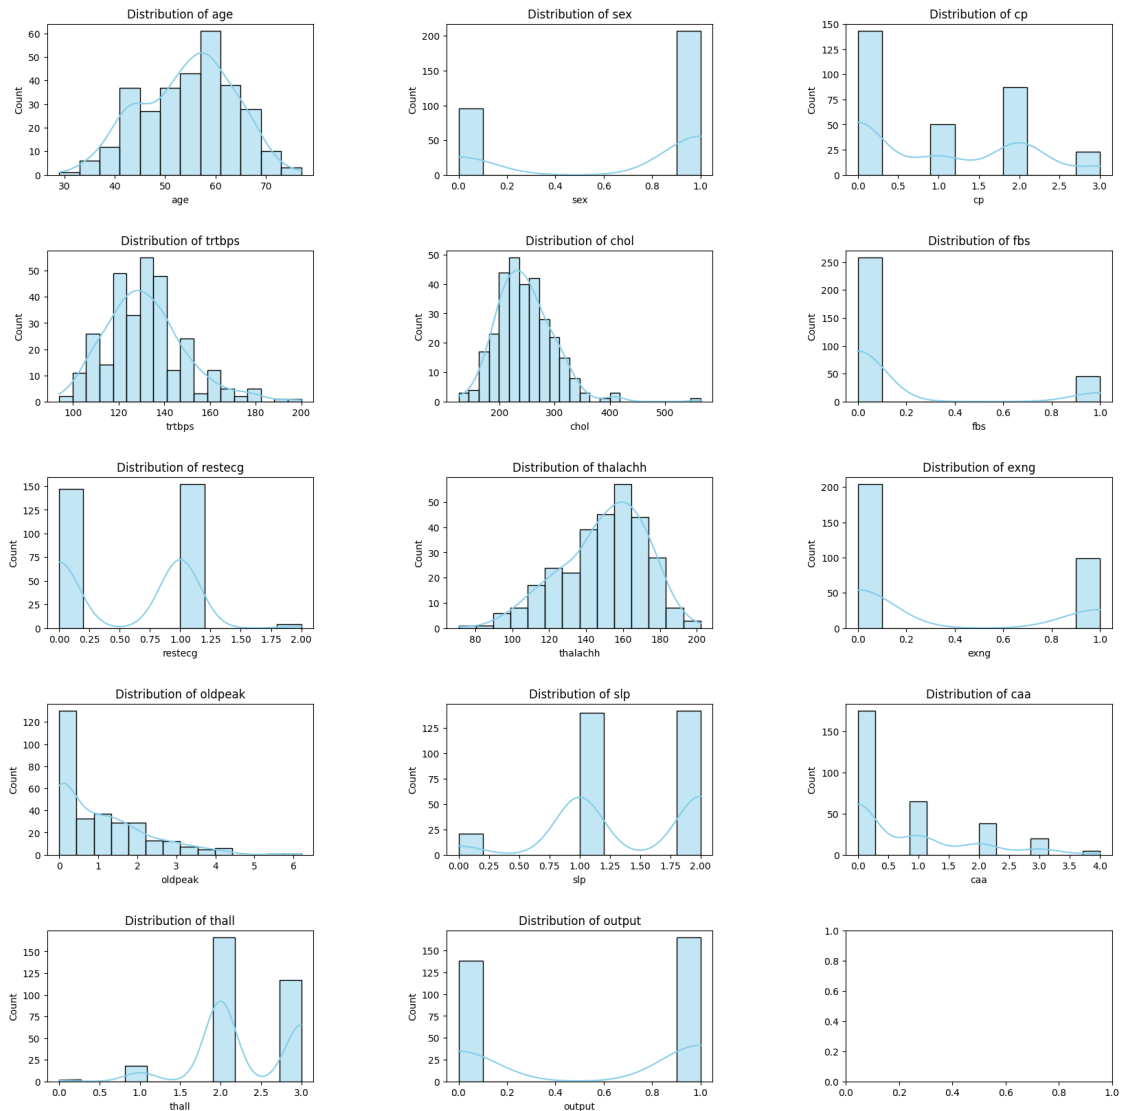
	thall	output
count	303.000000	303.000000
mean	2.313531	0.544554
std	0.612277	0.498835
min	0.000000	0.000000
25%	2.000000	0.000000
50%	2.000000	1.000000
75%	3.000000	1.000000
max	3.000000	1.000000

Visualizations

```
[6]: fig, axes = plt.subplots(nrows=5, ncols=3, figsize=(20, 20))
fig.subplots_adjust(hspace=0.5, wspace=0.5)

# Plot histograms for each attribute
for i, col in enumerate(df):
    sns.histplot(df[col], ax=axes[i//3, i%3], kde=True, color='skyblue').
    set(title=f'Distribution of {col}')

plt.show()
```



```
[7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 303 entries, 0 to 302
```

```
Data columns (total 14 columns):
```

#	Column	Non-Null Count	Dtype
0	age	303 non-null	int64
1	sex	303 non-null	int64
2	cp	303 non-null	int64
3	trtbps	303 non-null	int64
4	chol	303 non-null	int64
5	fbs	303 non-null	int64

```

6  restecg    303 non-null    int64
7  thalachh   303 non-null    int64
8  exng       303 non-null    int64
9  oldpeak    303 non-null    float64
10 slp        303 non-null    int64
11 caa        303 non-null    int64
12 thall      303 non-null    int64
13 output     303 non-null    int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB

```

```
[8]: df.isnull().sum()
```

```

[8]: age          0
sex            0
cp             0
trtbps         0
chol           0
fbs            0
restecg        0
thalachh       0
exng           0
oldpeak        0
slp            0
caa            0
thall          0
output         0
dtype: int64

```

There are no categorical attributes and there are no missing values, outliers Hence no transformations are required

2. Analyzing and discuss the relationships between the data attributes, and between the data attributes and label. This involves computing the Pearson Correlation Coefficient (PCC) and generating scatter plots.

Computing the PCC & Scatter Plots

```

[9]: correlation_matrix = df.corr()
print(correlation_matrix)
sns.pairplot(df)
plt.show()

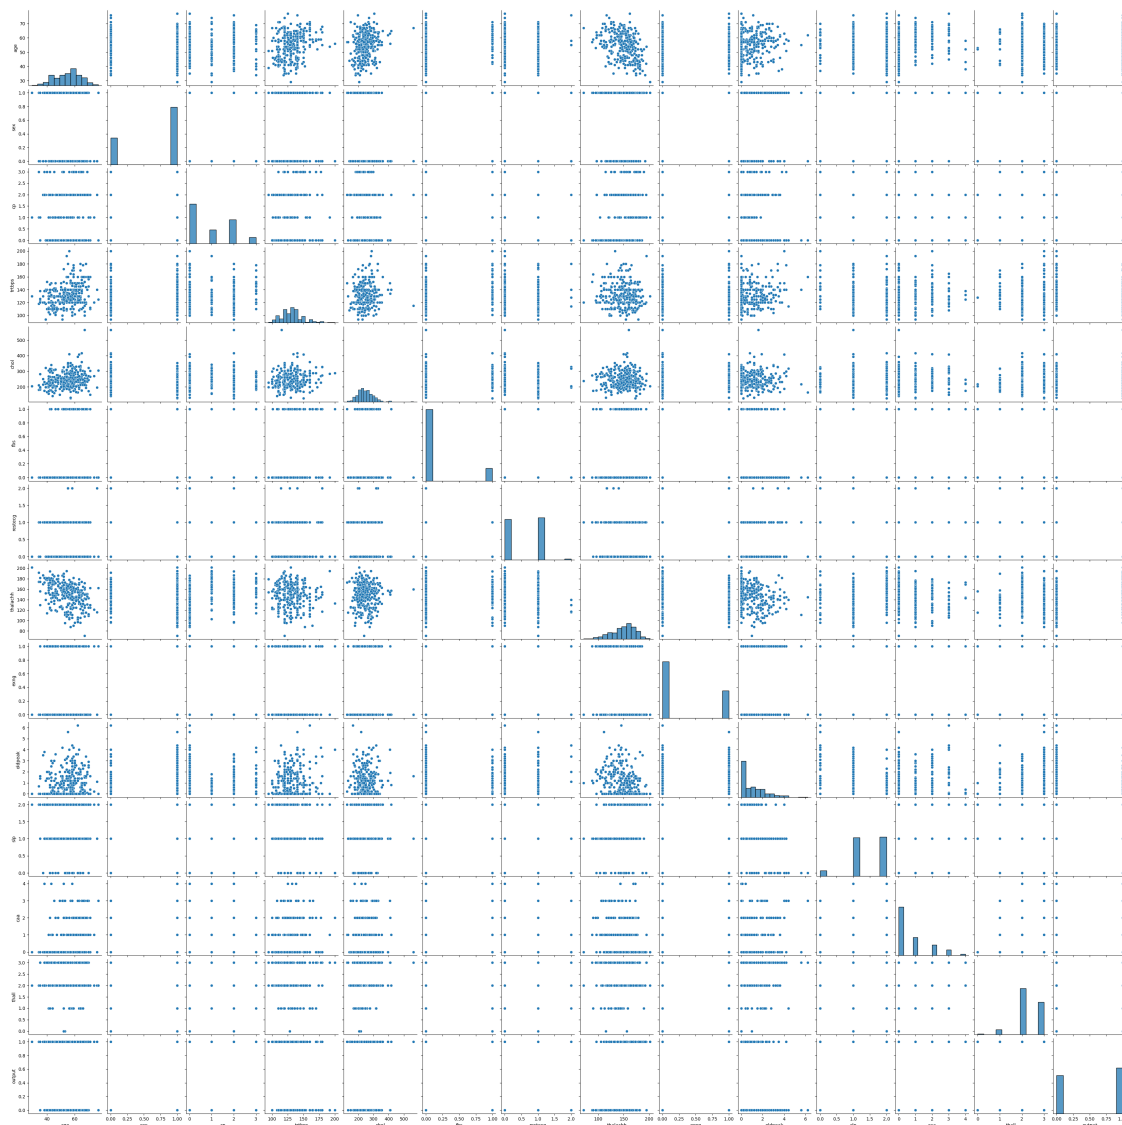
```

	age	sex	cp	trtbps	chol	fbs \
age	1.000000	-0.098447	-0.068653	0.279351	0.213678	0.121308
sex	-0.098447	1.000000	-0.049353	-0.056769	-0.197912	0.045032
cp	-0.068653	-0.049353	1.000000	0.047608	-0.076904	0.094444
trtbps	0.279351	-0.056769	0.047608	1.000000	0.123174	0.177531
chol	0.213678	-0.197912	-0.076904	0.123174	1.000000	0.013294

fbs	0.121308	0.045032	0.094444	0.177531	0.013294	1.000000
restecg	-0.116211	-0.058196	0.044421	-0.114103	-0.151040	-0.084189
thalachh	-0.398522	-0.044020	0.295762	-0.046698	-0.009940	-0.008567
exng	0.096801	0.141664	-0.394280	0.067616	0.067023	0.025665
oldpeak	0.210013	0.096093	-0.149230	0.193216	0.053952	0.005747
slp	-0.168814	-0.030711	0.119717	-0.121475	-0.004038	-0.059894
caa	0.276326	0.118261	-0.181053	0.101389	0.070511	0.137979
thall	0.068001	0.210041	-0.161736	0.062210	0.098803	-0.032019
output	-0.225439	-0.280937	0.433798	-0.144931	-0.085239	-0.028046

	restecg	thalachh	exng	oldpeak	slp	caa \
age	-0.116211	-0.398522	0.096801	0.210013	-0.168814	0.276326
sex	-0.058196	-0.044020	0.141664	0.096093	-0.030711	0.118261
cp	0.044421	0.295762	-0.394280	-0.149230	0.119717	-0.181053
trtbps	-0.114103	-0.046698	0.067616	0.193216	-0.121475	0.101389
chol	-0.151040	-0.009940	0.067023	0.053952	-0.004038	0.070511
fbs	-0.084189	-0.008567	0.025665	0.005747	-0.059894	0.137979
restecg	1.000000	0.044123	-0.070733	-0.058770	0.093045	-0.072042
thalachh	0.044123	1.000000	-0.378812	-0.344187	0.386784	-0.213177
exng	-0.070733	-0.378812	1.000000	0.288223	-0.257748	0.115739
oldpeak	-0.058770	-0.344187	0.288223	1.000000	-0.577537	0.222682
slp	0.093045	0.386784	-0.257748	-0.577537	1.000000	-0.080155
caa	-0.072042	-0.213177	0.115739	0.222682	-0.080155	1.000000
thall	-0.011981	-0.096439	0.206754	0.210244	-0.104764	0.151832
output	0.137230	0.421741	-0.436757	-0.430696	0.345877	-0.391724

	thall	output
age	0.068001	-0.225439
sex	0.210041	-0.280937
cp	-0.161736	0.433798
trtbps	0.062210	-0.144931
chol	0.098803	-0.085239
fbs	-0.032019	-0.028046
restecg	-0.011981	0.137230
thalachh	-0.096439	0.421741
exng	0.206754	-0.436757
oldpeak	0.210244	-0.430696
slp	-0.104764	0.345877
caa	0.151832	-0.391724
thall	1.000000	-0.344029
output	-0.344029	1.000000

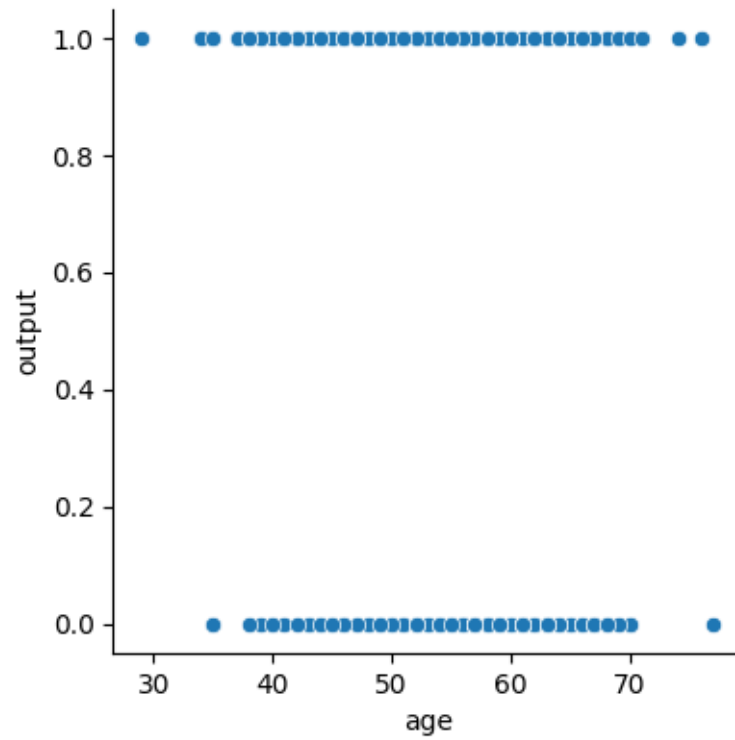


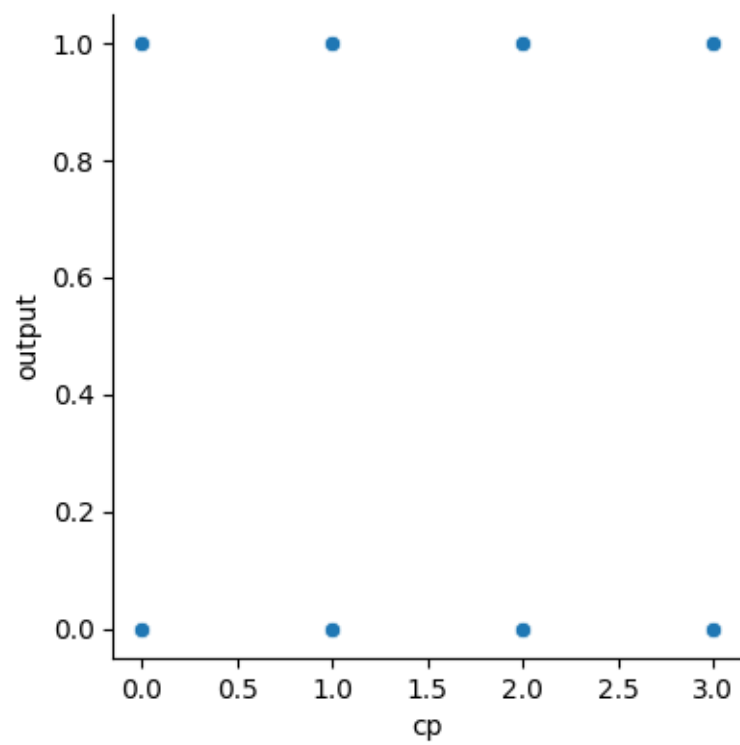
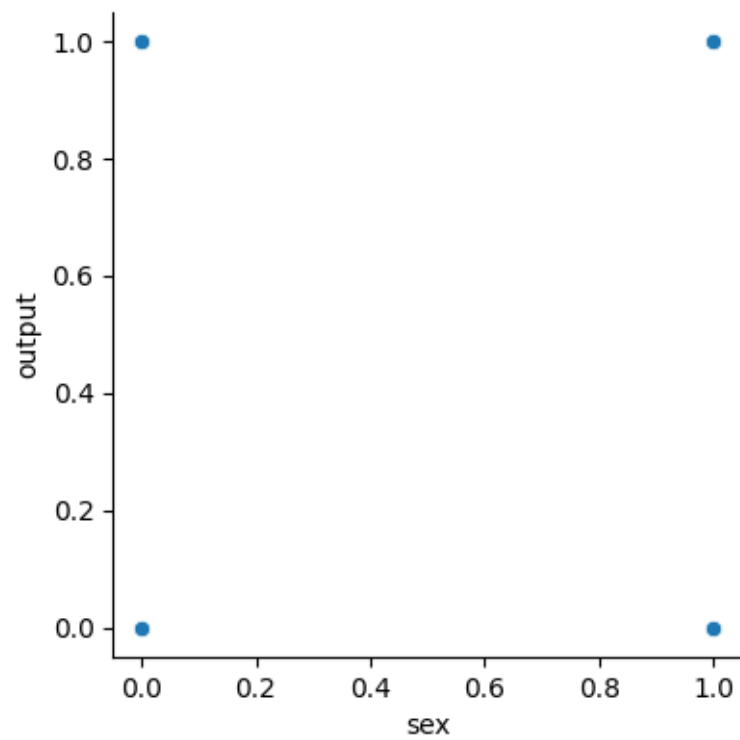
```
[10]: correlations_matrix_2 = (df).corr()
print(correlations_matrix_2["output"].sort_values(ascending=False))

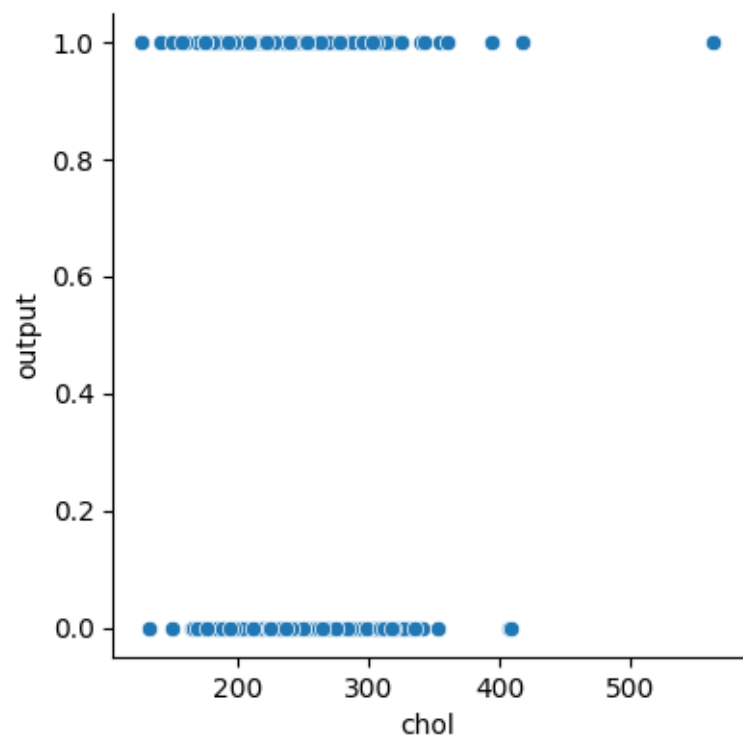
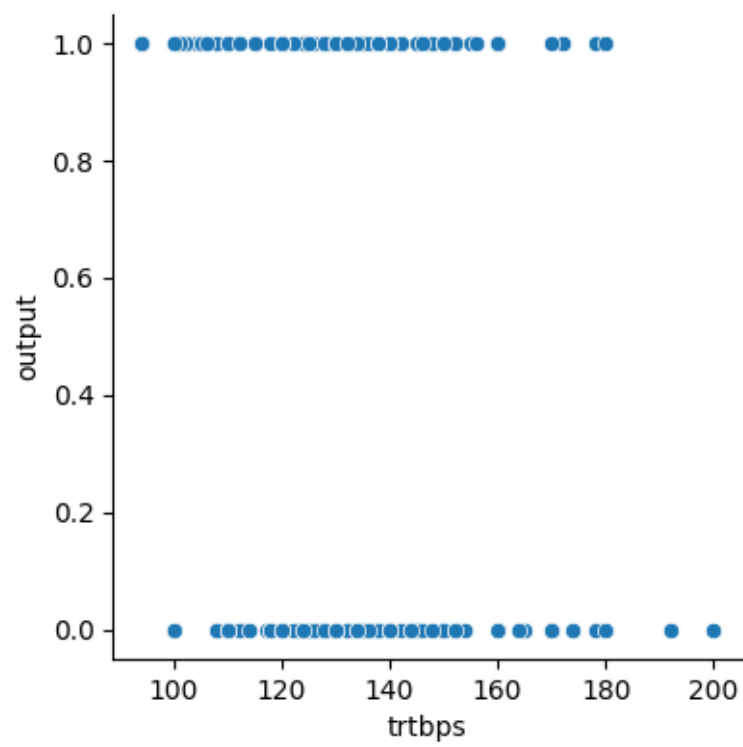
for attr in df.columns:
    sns.pairplot(data=df, x_vars=[attr], y_vars=["output"], kind="scatter",
    ↪height=4)
    plt.show()
```

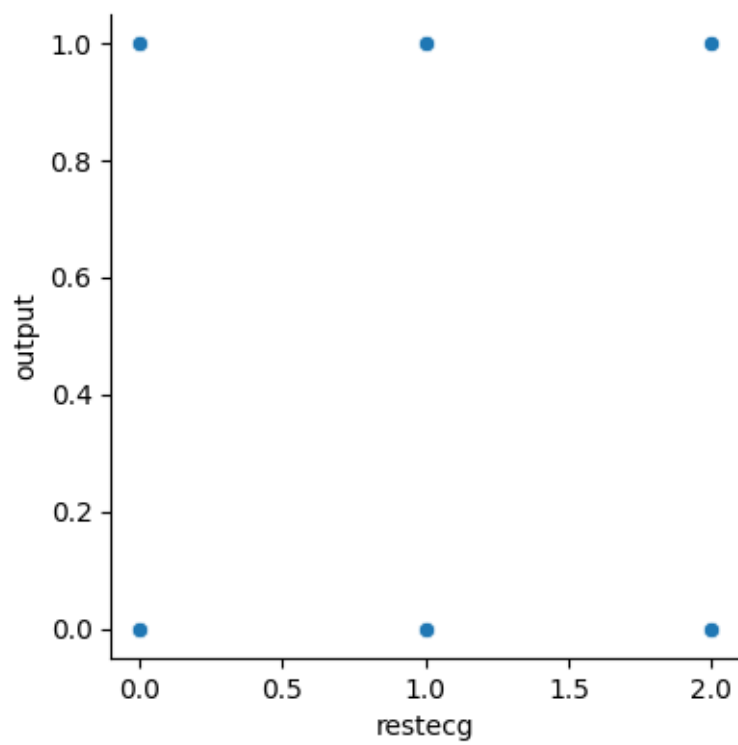
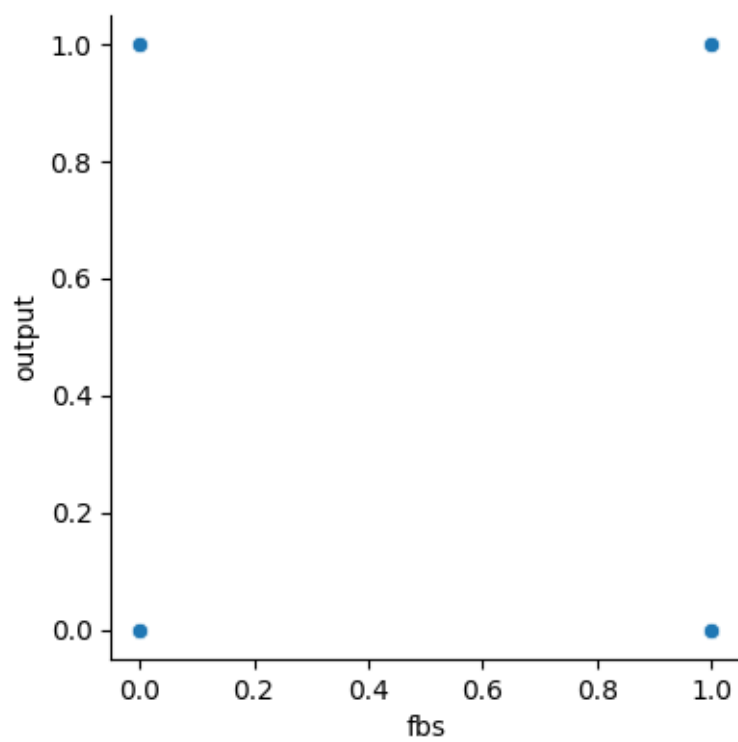
```
output      1.000000
cp           0.433798
thalachh     0.421741
slp          0.345877
```

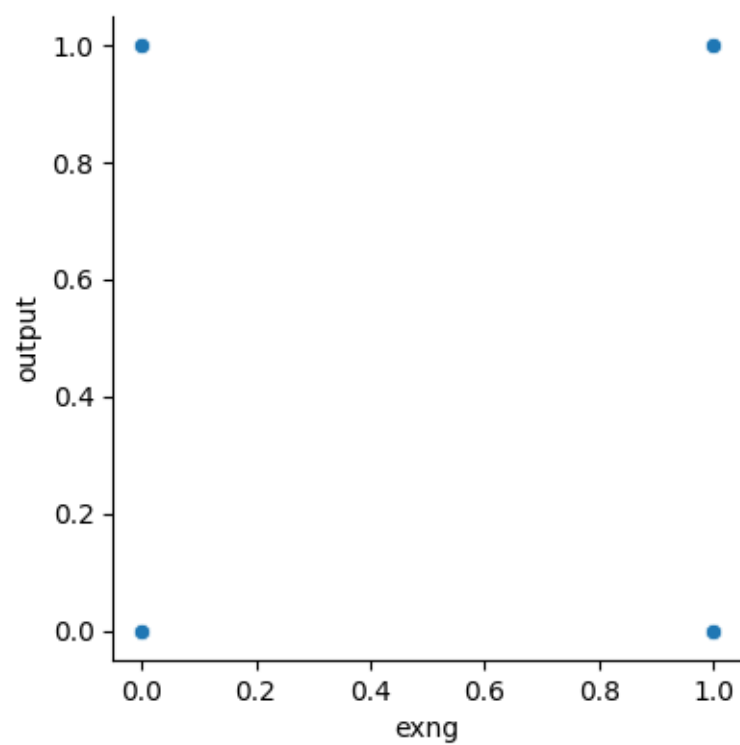
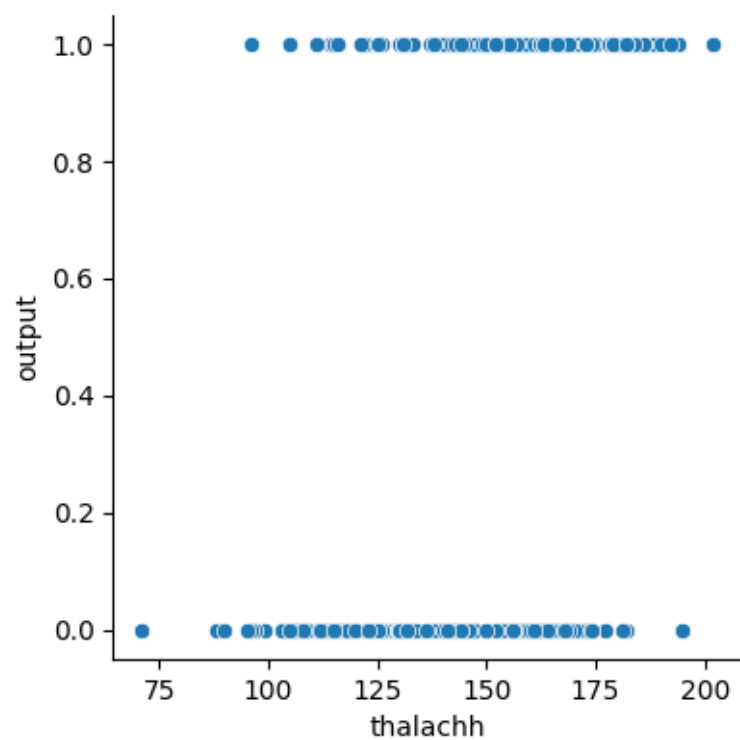
```
restecg    0.137230
fbs        -0.028046
chol       -0.085239
trtbps     -0.144931
age        -0.225439
sex        -0.280937
thall      -0.344029
caa        -0.391724
oldpeak    -0.430696
exng       -0.436757
Name: output, dtype: float64
```

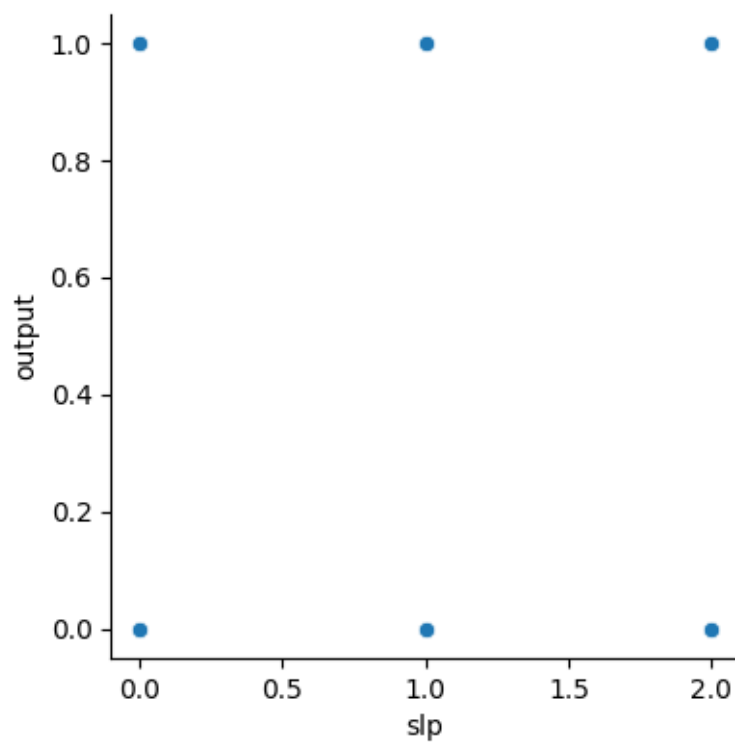
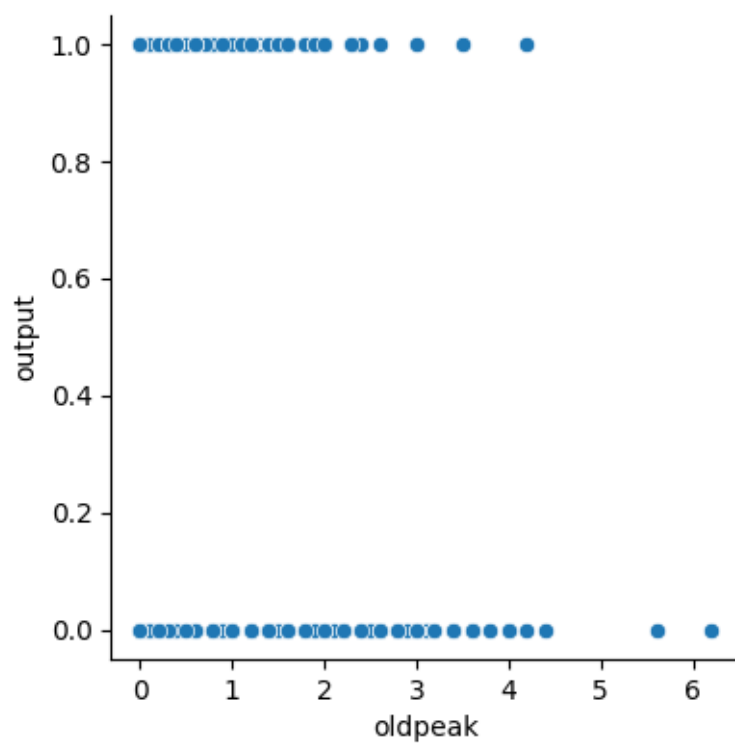


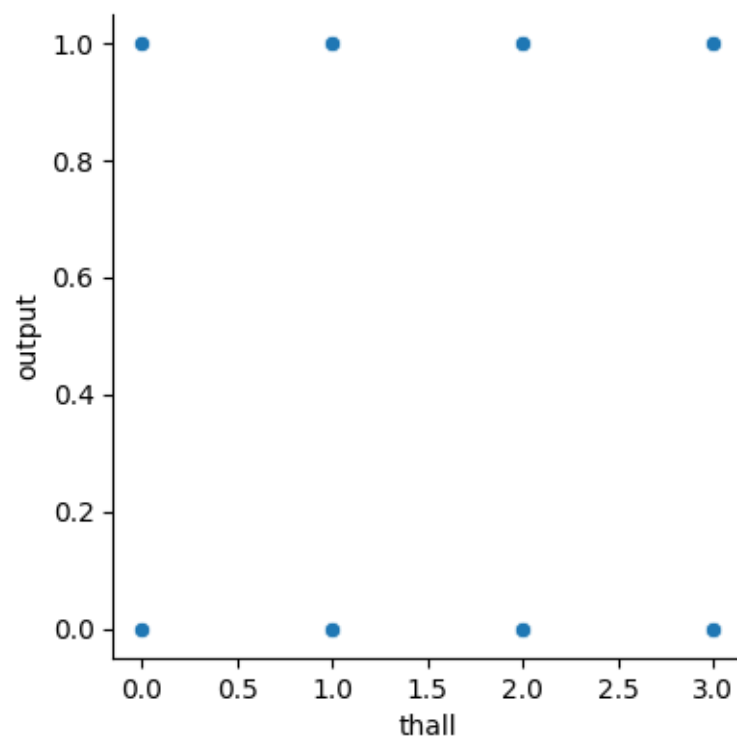
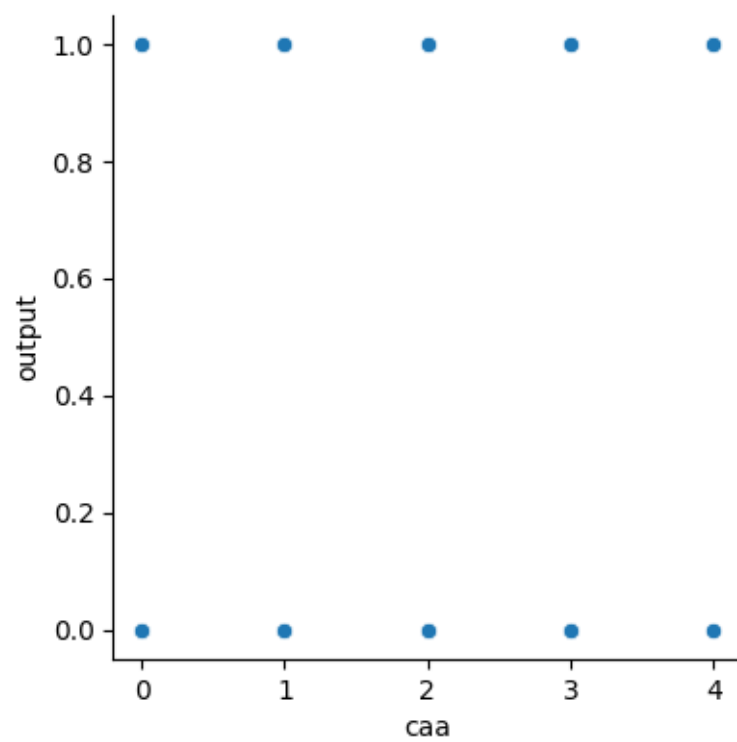


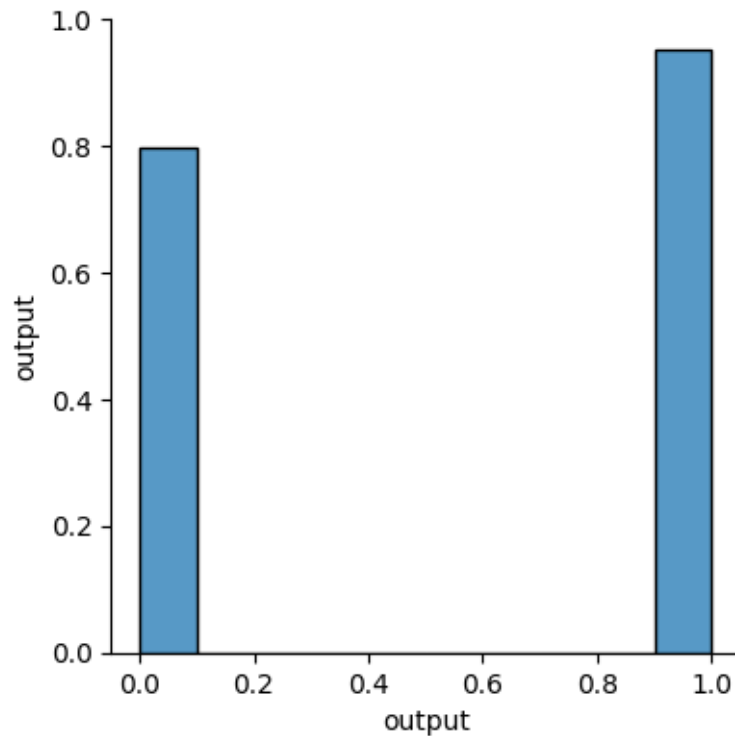












Positive Correlation (closer to 1 indicates a stronger positive relationship):

cp (chest pain type) and thalachh (maximum heart rate achieved) show moderate positive correlations with the output, indicating that higher values of these attributes are associated with a higher likelihood of the outcome. slp (slope of the peak exercise ST segment) and restecg (resting electrocardiographic results) show weaker positive correlations, suggesting a lesser but still positive association with the outcome.

Negative Correlation (closer to -1 indicates a stronger negative relationship):

exng (exercise induced angina), oldpeak (ST depression induced by exercise relative to rest), caa (number of major vessels colored by fluoroscopy), and thall (thalassemia) are negatively correlated with output, indicating that higher values are associated with a lower likelihood of the outcome. sex (gender), age, trtbps (resting blood pressure), chol (serum cholesterol), and fbs (fasting blood sugar) also show negative correlations with varying strengths, suggesting that higher values may decrease the likelihood of the outcome, though to a lesser extent for some of these factors.

fbs(Fasting blood sugar), chol(Cholestrol) exhibit minimal correlations with the label.

```
[11]: # Drop the specified columns from the DataFrame
df_updated=df.copy()
df_updated.drop(columns=['fbs','chol'],axis=1, inplace=True)
```

PreProcessing

```
[12]: attributes = ['age', 'sex', 'cp', 'trtbps', 'restecg',
                  'thalachh', 'exng', 'oldpeak', 'slp', 'caa', 'thall']

# Creating a pipeline for numerical data
num_con_pipeline = make_pipeline(StandardScaler())

# Applying ColumnTransformer to the specified attributes
prep = ColumnTransformer([
    ("cont", num_con_pipeline, attributes)
])

# Assuming 'df' is your DataFrame and you've separated features and label
↳('output')
X = df_updated.drop('output', axis=1)
y = df['output']

# Transforming the features
X_transformed = prep.fit_transform(X)

# Converting transformed features back to DataFrame with meaningful column names
X_transformed_df = pd.DataFrame(X_transformed, columns=prep.
    ↳get_feature_names_out(), index=X.index)

# Now, X_transformed_df contains the standardized features, ready for modeling.
```

```
[13]: X_transformed_df
```

```
[13]:
```

	cont__age	cont__sex	cont__cp	cont__trtbps	cont__restecg	\
0	0.952197	0.681005	1.973123	0.763956	-1.005832	
1	-1.915313	0.681005	1.002577	-0.092738	0.898962	
2	-1.474158	-1.468418	0.032031	-0.092738	-1.005832	
3	0.180175	0.681005	0.032031	-0.663867	0.898962	
4	0.290464	-1.468418	-0.938515	-0.663867	0.898962	
..	
298	0.290464	-1.468418	-0.938515	0.478391	0.898962	
299	-1.033002	0.681005	1.973123	-1.234996	0.898962	
300	1.503641	0.681005	-0.938515	0.706843	0.898962	
301	0.290464	0.681005	-0.938515	-0.092738	0.898962	
302	0.290464	-1.468418	0.032031	-0.092738	-1.005832	
	cont__thalachh	cont__exng	cont__oldpeak	cont__slp	cont__caa	\
0	0.015443	-0.696631	1.087338	-2.274579	-0.714429	
1	1.633471	-0.696631	2.122573	-2.274579	-0.714429	
2	0.977514	-0.696631	0.310912	0.976352	-0.714429	
3	1.239897	-0.696631	-0.206705	0.976352	-0.714429	

4	0.583939	1.435481	-0.379244	0.976352	-0.714429
..
298	-1.165281	1.435481	-0.724323	-0.649113	-0.714429
299	-0.771706	-0.696631	0.138373	-0.649113	-0.714429
300	-0.378132	-0.696631	2.036303	-0.649113	1.244593
301	-1.515125	1.435481	0.138373	-0.649113	0.265082
302	1.064975	-0.696631	-0.896862	-0.649113	0.265082

	cont__thall
0	-2.148873
1	-0.512922
2	-0.512922
3	-0.512922
4	-0.512922
..	...
298	1.123029
299	1.123029
300	1.123029
301	1.123029
302	-0.512922

[303 rows x 11 columns]

3. Splitting training data, for validation, and for testing.

```
[14]: X = X_transformed_df
# First split: Separate out the training set
X_train, X_temp, y_train, y_temp = train_test_split(X, y, test_size=0.4,
↳ random_state=42)

# Second split: Divide the remaining data into validation and test sets
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5,
↳ random_state=42)

print("Training set size:", X_train.shape[0])
print("Validation set size:", X_val.shape[0])
print("Test set size:", X_test.shape[0])
```

Training set size: 181
Validation set size: 61
Test set size: 61

Verification of splitting

```
[15]: X_train.describe()
```

```
[15]:
```

	cont__age	cont__sex	cont__cp	cont__trtbps	cont__restecg	\
count	181.000000	181.000000	181.000000	181.000000	181.000000	

mean	0.039419	-0.055261	0.074928	0.014231	-0.037649
std	0.999323	1.022859	1.022140	0.977338	0.975788
min	-2.246179	-1.468418	-0.938515	-2.148802	-1.005832
25%	-0.591847	-1.468418	-0.938515	-0.663867	-1.005832
50%	0.180175	0.681005	0.032031	-0.092738	0.898962
75%	0.731619	0.681005	1.002577	0.478391	0.898962
max	2.496240	0.681005	1.973123	3.448262	2.803756

	cont__thalachh	cont__exng	cont__oldpeak	cont__slp	cont__caa \
count	181.000000	181.000000	181.000000	181.000000	181.000000
mean	-0.012825	0.033706	-0.017008	-0.020480	-0.086676
std	1.024189	1.014621	0.932945	1.005865	0.910873
min	-2.695849	-0.696631	-0.896862	-2.274579	-0.714429
25%	-0.771706	-0.696631	-0.896862	-0.649113	-0.714429
50%	0.190365	-0.696631	-0.206705	-0.649113	-0.714429
75%	0.802592	1.435481	0.483451	0.976352	0.265082
max	1.983316	1.435481	3.934233	0.976352	2.224104

	cont__thall
count	181.000000
mean	0.020344
std	0.988874
min	-3.784824
25%	-0.512922
50%	-0.512922
75%	1.123029
max	1.123029

```
[16]: X_test.describe()
```

```
[16]:
```

	cont__age	cont__sex	cont__cp	cont__trtbps	cont__restecg \
count	61.000000	61.000000	61.000000	61.000000	61.000000
mean	-0.004243	-0.023724	0.032031	-0.012218	-0.069048
std	0.980624	1.017392	0.954233	1.150480	1.078802
min	-2.246179	-1.468418	-0.938515	-2.148802	-1.005832
25%	-0.922713	-1.468418	-0.938515	-0.663867	-1.005832
50%	0.180175	0.681005	0.032031	-0.206964	-1.005832
75%	0.621330	0.681005	1.002577	0.478391	0.898962
max	2.385951	0.681005	1.973123	3.905165	2.803756

	cont__thalachh	cont__exng	cont__oldpeak	cont__slp	cont__caa \
count	61.000000	61.000000	61.000000	61.000000	61.000000
mean	0.039817	-0.102435	-0.003053	-0.036233	-0.056069
std	1.023285	0.963874	1.067448	1.035033	1.007220
min	-3.439267	-0.696631	-0.896862	-2.274579	-0.714429
25%	-0.509323	-0.696631	-0.896862	-0.649113	-0.714429
50%	0.190365	-0.696631	-0.292975	-0.649113	-0.714429

75%	0.715131	1.435481	0.483451	0.976352	0.265082
max	1.852124	1.435481	4.451851	0.976352	3.203615

	cont__thall
count	61.000000
mean	-0.030182
std	0.913024
min	-2.148873
25%	-0.512922
50%	-0.512922
75%	1.123029
max	1.123029

```
[17]: X_val.describe()
```

```
[17]:
```

	cont__age	cont__sex	cont__cp	cont__trtbps	cont__restecg	\
count	61.000000	61.000000	61.000000	61.000000	61.000000	
mean	-0.112723	0.187695	-0.254359	-0.030007	0.180761	
std	1.036732	0.911369	0.959880	0.926902	0.993659	
min	-2.797624	-1.468418	-0.938515	-1.806125	-1.005832	
25%	-0.922713	0.681005	-0.938515	-0.663867	-1.005832	
50%	-0.040403	0.681005	-0.938515	-0.092738	0.898962	
75%	0.511041	0.681005	0.032031	0.478391	0.898962	
max	1.834507	0.681005	1.973123	2.191778	2.803756	

	cont__thalachh	cont__exng	cont__oldpeak	cont__slp	cont__caa	\
count	61.000000	61.000000	61.000000	61.000000	61.000000	
mean	-0.001763	0.002422	0.053518	0.097002	0.313255	
std	0.924059	1.009198	1.138245	0.965010	1.191979	
min	-2.214813	-0.696631	-0.896862	-2.274579	-0.714429	
25%	-0.421862	-0.696631	-0.896862	-0.649113	-0.714429	
50%	0.059173	-0.696631	-0.551783	0.976352	0.265082	
75%	0.583939	1.435481	0.655990	0.976352	1.244593	
max	2.289429	1.435481	2.898999	0.976352	3.203615	

	cont__thall
count	61.000000
mean	-0.030182
std	1.131224
min	-3.784824
25%	-0.512922
50%	-0.512922
75%	1.123029
max	1.123029

If we look at the data given by test and validate describe , the mean , median and standard deviation and the quartile range looks similiar. This means test, validate portion of the data is the

representative of the entire dataset.

The consistency across training, validation, and test sets regarding key statistical measures supports the validity of your data splitting strategy. It suggests that any conclusions drawn from the model's performance on the validation and test sets should be applicable to the entire dataset.

4. Train different classifiers and tweak the hyperparameters to improve performance (use the grid search if you want or manually try different values). Reporting training, validation and testing performance (classification accuracy, precision, recall and F1 score) and discuss the impact of the hyperparameters

A. Multinomial Logistic Regression (softmax regression); hyperparameters to explore: C, solver, max number of iterations. [10 points]

```
[18]: from sklearn.linear_model import LogisticRegression as lg
      from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
      from sklearn.model_selection import GridSearchCV
      import warnings

      warnings.filterwarnings('ignore')

      # Hyperparameter grid
      param_grid = {'C': [0.001, 0.01, 0.1, 1, 10, 100],
                    'solver': ['newton-cg', 'lbfgs', 'liblinear'],
                    'max_iter': [100, 200, 300, 400, 500]}

      sfmax_reg = lg(multi_class="multinomial", solver="lbfgs", C=10)

      grid = GridSearchCV(sfmax_reg, param_grid, cv=3, scoring='accuracy')
      grid.fit(X_train, y_train)

      best_params = grid.best_params_

      best_sf_train = lg(multi_class="multinomial", **best_params)
      best_sf_train.fit(X_train, y_train)

      # Make predictions on the training, validation, and test sets
      train_predictions = best_sf_train.predict(X_train)
      val_predictions = best_sf_train.predict(X_val)
      test_predictions = best_sf_train.predict(X_test)

      # Calculate accuracy for training, validation, and test sets
      train_accuracy = accuracy_score(y_train, train_predictions)
      val_accuracy = accuracy_score(y_val, val_predictions)
      test_accuracy = accuracy_score(y_test, test_predictions)

      # Calculate F1 scores, precision, and recall for training set
```

```

precision_train = precision_score(y_train, train_predictions,
    ↪average='weighted')
recall_train = recall_score(y_train, train_predictions, average='weighted')
f1_train = f1_score(y_train, train_predictions, average='weighted')

# Calculate F1 scores, precision, and recall for validation set
precision_val = precision_score(y_val, val_predictions, average='weighted')
recall_val = recall_score(y_val, val_predictions, average='weighted')
f1_val = f1_score(y_val, val_predictions, average='weighted')

# Calculate F1 scores, precision, and recall for test set
precision_test = precision_score(y_test, test_predictions, average='weighted')
recall_test = recall_score(y_test, test_predictions, average='weighted')
f1_test = f1_score(y_test, test_predictions, average='weighted')

# Print results for Logistic Regression model
print("Logistic Regression Model Results:")
print("\nTraining Set:")
print(f"Accuracy: {train_accuracy}")
print(f"Precision: {precision_train}")
print(f"Recall: {recall_train}")
print(f"F1 Score: {f1_train}")

print("\nValidation Set:")
print(f"Accuracy: {val_accuracy}")
print(f"Precision: {precision_val}")
print(f"Recall: {recall_val}")
print(f"F1 Score: {f1_val}")

print("\nTest Set:")
print(f"Accuracy: {test_accuracy}")
print(f"Precision: {precision_test}")
print(f"Recall: {recall_test}")
print(f"F1 Score: {f1_test}")

print(f"\nBest parameters: {best_params}")

```

Logistic Regression Model Results:

Training Set:

Accuracy: 0.8839779005524862

Precision: 0.8872750177998796

Recall: 0.8839779005524862

F1 Score: 0.8833984517714492

Validation Set:

Accuracy: 0.8688524590163934

Precision: 0.8711440155120748
Recall: 0.8688524590163934
F1 Score: 0.8690644431882419

Test Set:
Accuracy: 0.8032786885245902
Precision: 0.8032786885245902
Recall: 0.8032786885245902
F1 Score: 0.8032786885245902

Best parameters: {'C': 1, 'max_iter': 100, 'solver': 'newton-cg'}

C (Regularization Parameter): Your observations indicate that the grid search found the optimal C value to be 1, exploring a range from 0.001 to 100. This optimal value points towards the effectiveness of moderate regularization in enhancing the model's performance. Moderate regularization helps in avoiding overfitting while still allowing the model to learn complex patterns, striking a balance between bias and variance.

Solver: Among the solvers 'newton-cg', 'lbfgs', and 'liblinear' considered during the grid search, 'newton-cg' emerged as the best solver. This indicates that for this particular logistic regression model, especially in the context of multiclass softmax regression, the Newton-Conjugate Gradient solver optimizes the cost function most efficiently. The choice of 'newton-cg' suggests that it effectively deals with the logistic regression's mathematical characteristics and data structure, contributing to higher cross-validated accuracy.

Max_iter (Maximum Number of Iterations): The grid search tested values from 100 to 500 for max_iter, with 100 being identified as the best-performing value. This outcome implies that the optimization algorithm converges swiftly, within the first 100 iterations, which is advantageous for computational efficiency. Early convergence without sacrificing accuracy indicates that the initial parameters are adequately close to the optimal solution or that the model is well-structured to reach an optimal solution quickly.

Summary: The identified optimal hyperparameters—C=1 for regularization strength, 'newton-cg' as the solver, and 100 maximum iterations—demonstrate a well-tuned logistic regression model for your dataset. These parameters collectively contribute to the model's strong performance by ensuring a good fit to the data, efficient optimization, and prevention of overfitting. The choice of a moderate regularization level, combined with an effective solver and a relatively low number of iterations, reflects a strategic approach to balancing model complexity, convergence speed, and regularization to achieve optimal performance in predicting the outcomes of your multiclass softmax regression model.

B. Support vector machines (make sure to try using kernels); hyperparameters : C, kernel, degree of polynomial kernel, gamma.

```
[19]: from sklearn.svm import SVC

param_grid_svm = {'C': [0.001, 0.01, 0.1, 1, 10, 100],
                  'kernel': ['linear', 'rbf', 'poly'],
                  'degree': [2, 3, 4],
                  'gamma': ['scale', 'auto']}
```

```

svm_model = SVC()

grid_search_svm = GridSearchCV(svm_model, param_grid_svm, cv=3,
    ↪scoring='accuracy')
grid_search_svm.fit(X_train, y_train)

best_params_svm = grid_search_svm.best_params_

best_svm = SVC(**best_params_svm)
best_svm.fit(X_train, y_train)

train_predictions_svm = best_svm.predict(X_train)
val_predictions_svm = best_svm.predict(X_val)
test_predictions_svm = best_svm.predict(X_test)

# Calculate accuracy for SVM model
train_accuracy_svm = accuracy_score(y_train, train_predictions_svm)
val_accuracy_svm = accuracy_score(y_val, val_predictions_svm)
test_accuracy_svm = accuracy_score(y_test, test_predictions_svm)

# Calculate F1 scores, precision, and recall for SVM model on training set
precision_train_svm = precision_score(y_train, train_predictions_svm,
    ↪average='weighted')
recall_train_svm = recall_score(y_train, train_predictions_svm,
    ↪average='weighted')
f1_train_svm = f1_score(y_train, train_predictions_svm, average='weighted')

# Calculate F1 scores, precision, and recall for SVM model on validation set
precision_val_svm = precision_score(y_val, val_predictions_svm,
    ↪average='weighted')
recall_val_svm = recall_score(y_val, val_predictions_svm, average='weighted')
f1_val_svm = f1_score(y_val, val_predictions_svm, average='weighted')

# Calculate F1 scores, precision, and recall for SVM model on test set
precision_test_svm = precision_score(y_test, test_predictions_svm,
    ↪average='weighted')
recall_test_svm = recall_score(y_test, test_predictions_svm, average='weighted')
f1_test_svm = f1_score(y_test, test_predictions_svm, average='weighted')

# Print results for SVM model
print("SVM Model Results:")
print("\nTraining Set:")
print(f"Accuracy: {train_accuracy_svm}")
print(f"Precision: {precision_train_svm}")
print(f"Recall: {recall_train_svm}")
print(f"F1 Score: {f1_train_svm}")

```

```

print("\nValidation Set:")
print(f"Accuracy: {val_accuracy_svm}")
print(f"Precision: {precision_val_svm}")
print(f"Recall: {recall_val_svm}")
print(f"F1 Score: {f1_val_svm}")

print("\nTest Set:")
print(f"Accuracy: {test_accuracy_svm}")
print(f"Precision: {precision_test_svm}")
print(f"Recall: {recall_test_svm}")
print(f"F1 Score: {f1_test_svm}")

print("\nBest Hyperparameters:")
print(best_params_svm)

```

SVM Model Results:

Training Set:

Accuracy: 0.8729281767955801
Precision: 0.8778129886375483
Recall: 0.8729281767955801
F1 Score: 0.8720643035807545

Validation Set:

Accuracy: 0.8688524590163934
Precision: 0.8688524590163934
Recall: 0.8688524590163934
F1 Score: 0.8688524590163934

Test Set:

Accuracy: 0.8360655737704918
Precision: 0.8351484580992778
Recall: 0.8360655737704918
F1 Score: 0.8346588138200098

Best Hyperparameters:

```
{'C': 0.1, 'degree': 2, 'gamma': 'scale', 'kernel': 'linear'}
```

C (Regularization Parameter): The best C value is 0.1. This suggests that a lower degree of regularization (since C is the inverse of regularization strength) helps to achieve better performance, likely by allowing the model to fit the training data more closely without significant overfitting.

Kernel: The optimal kernel is 'linear', which implies that the decision boundary between the classes in your dataset can be well approximated using a linear function. This could indicate that the feature space is linearly separable or close to it.

Degree: The best degree for polynomial kernels is 2, but since the best kernel is linear, this parameter does not affect the model.

Gamma: The 'scale' option for gamma is chosen, which is typically effective for features of varying scales and distributions. This choice automatically adjusts gamma based on the feature variance, offering a balanced approach to handling different feature characteristics.

In summary, the identification of a linear kernel and a C value of 0.1 as optimal hyperparameters demonstrates a well-tuned approach to SVM modeling for dataset. These hyperparameter choices indicate a model that balances complexity with the ability to generalize, leading to robust performance across different subsets of the data.

C. Random Forest classifier (also analyze feature importance); hyperparameters: the number of trees, max depth, the minimum number of samples required to split an internal node, the minimum number of samples required to be at a leaf node.

```
[20]: from sklearn.ensemble import RandomForestClassifier

param_grid_rf = {'n_estimators': [50, 100, 150],
                  'max_depth': [None, 10, 20, 30],
                  'min_samples_split': [2, 5, 10],
                  'min_samples_leaf': [1, 2, 4]}

rf_model = RandomForestClassifier()

grid_search_rf = GridSearchCV(rf_model, param_grid_rf, cv=3, scoring='accuracy')
grid_search_rf.fit(X_train, y_train)

best_params_rf = grid_search_rf.best_params_

best_rf = RandomForestClassifier(**best_params_rf)
best_rf.fit(X_train, y_train)

train_predictions_rf = best_rf.predict(X_train)
val_predictions_rf = best_rf.predict(X_val)
test_predictions_rf = best_rf.predict(X_test)

# Calculate accuracy for RandomForest model
train_accuracy_rf = accuracy_score(y_train, train_predictions_rf)
val_accuracy_rf = accuracy_score(y_val, val_predictions_rf)
test_accuracy_rf = accuracy_score(y_test, test_predictions_rf)

# Calculate F1 scores, precision, and recall for RandomForest model on training
↪set
precision_train_rf = precision_score(y_train, train_predictions_rf,
↪average='weighted')
recall_train_rf = recall_score(y_train, train_predictions_rf,
↪average='weighted')
f1_train_rf = f1_score(y_train, train_predictions_rf, average='weighted')
```



```

# Calculate F1 scores, precision, and recall for RandomForest model on
    validation set
precision_val_rf = precision_score(y_val, val_predictions_rf,
    average='weighted')
recall_val_rf = recall_score(y_val, val_predictions_rf, average='weighted')
f1_val_rf = f1_score(y_val, val_predictions_rf, average='weighted')

# Calculate F1 scores, precision, and recall for RandomForest model on test set
precision_test_rf = precision_score(y_test, test_predictions_rf,
    average='weighted')
recall_test_rf = recall_score(y_test, test_predictions_rf, average='weighted')
f1_test_rf = f1_score(y_test, test_predictions_rf, average='weighted')

# Print results for RandomForest model
print("RandomForest Model Results:")
print("\nTraining Set:")
print(f"Accuracy: {train_accuracy_rf}")
print(f"Precision: {precision_train_rf}")
print(f"Recall: {recall_train_rf}")
print(f"F1 Score: {f1_train_rf}")

print("\nValidation Set:")
print(f"Accuracy: {val_accuracy_rf}")
print(f"Precision: {precision_val_rf}")
print(f"Recall: {recall_val_rf}")
print(f"F1 Score: {f1_val_rf}")

print("\nTest Set:")
print(f"Accuracy: {test_accuracy_rf}")
print(f"Precision: {precision_test_rf}")
print(f"Recall: {recall_test_rf}")
print(f"F1 Score: {f1_test_rf}")

print("\nBest Hyperparameters:")
print(best_params_rf)

```

RandomForest Model Results:

Training Set:

Accuracy: 0.9392265193370166

Precision: 0.9396116430996282

Recall: 0.9392265193370166

F1 Score: 0.9391594424965998

Validation Set:

Accuracy: 0.8360655737704918

Precision: 0.8436263425664218

Recall: 0.8360655737704918
F1 Score: 0.8362418473470827

Test Set:
Accuracy: 0.819672131147541
Precision: 0.8186967775818734
Recall: 0.819672131147541
F1 Score: 0.8189559353563539

Best Hyperparameters:
{'max_depth': 30, 'min_samples_leaf': 4, 'min_samples_split': 5, 'n_estimators': 50}

Max_depth: The optimal configuration having max_depth as None allows trees to expand until all leaves are pure or contain less than min_samples_split samples. This setting is typically responsible for the model's high training performance but requires careful monitoring to prevent overfitting.

Min_samples_leaf: A min_samples_leaf of 2 means a split point at any depth will only be considered if it leaves at least two training samples in each of the left and right branches. This helps in making the model more general and less likely to overfit.

Min_samples_split: With min_samples_split set to 2, the smallest split includes only two samples. This allows the model to learn detailed patterns but, combined with min_samples_leaf, ensures a balance to prevent too fine-grained learning.

N_estimators: The chosen number of trees, 50, suggests that adding more trees beyond this point might not significantly improve the model's performance on this dataset. It's a balance between computational efficiency and model accuracy.

5. Combine your classifiers into an ensemble and try to outperform each individual classifier on the validation set. Once you have found a good one, try it on the test set. Describe and discuss your findings. [8 points]

Hard Voting

```
[21]: from sklearn.ensemble import VotingClassifier
      from sklearn.metrics import accuracy_score

      hard_voting_clf = VotingClassifier(
          estimators=[('lr', best_sf_train), ('svm', best_svm), ('rf', best_rf)],
          voting='hard'
      )

      hard_voting_clf.fit(X_train, y_train)

      train_accuracy_hard = accuracy_score(y_train, hard_voting_clf.predict(X_train))
      val_accuracy_hard = accuracy_score(y_val, hard_voting_clf.predict(X_val))

      print("Training Accuracy:", train_accuracy_hard)
      print("Validation Accuracy:", val_accuracy_hard)
```

```
#print("Testing Accuracy:", test_accuracy_hard)
```

Training Accuracy: 0.8895027624309392
Validation Accuracy: 0.8688524590163934

Soft Voting

```
[22]: from sklearn.ensemble import VotingClassifier
      from sklearn.metrics import accuracy_score

      best_svm = SVC(**best_params_svm, probability=True)

      soft_voting_clf = VotingClassifier(
          estimators=[('lr', best_sf_train), ('svm', best_svm), ('rf', best_rf)],
          voting='soft'
      )

      soft_voting_clf.fit(X_train, y_train)

      train_accuracy_soft = accuracy_score(y_train, soft_voting_clf.predict(X_train))
      val_accuracy_soft = accuracy_score(y_val, soft_voting_clf.predict(X_val))

      print("Training Accuracy:", train_accuracy_soft)
      print("Validation Accuracy:", val_accuracy_soft)
      #print("Test Accuracy:", test_accuracy_soft)
```

Training Accuracy: 0.8895027624309392
Validation Accuracy: 0.8524590163934426

It appears that assembling the previously trained classifiers into an ensemble using hard voting reduced the score on the validation set compared to our other classifiers. The hard voting ensemble gave a validation accuracy of 0.86. So, Using Hard voting ensemble to calculate the Test data Accuracy

```
[25]: test_accuracy_hard = accuracy_score(y_test, hard_voting_clf.predict(X_test))
      print("Testing Accuracy:", test_accuracy_hard)
```

Testing Accuracy: 0.8360655737704918